



# Apache Spark Programming

## With Databricks



# Welcome

2



# Course Agenda

Here's where we're headed

<b>DataFrames</b>	Introduction	Databricks Platform	Spark SQL	Reader & Writer	DataFrame & Column
<b>Transformations</b>	Aggregation	Datetimes	Complex Types	Additional Functions	User-Defined Functions
<b>Performance</b>	Spark Architecture	Query Optimization	Partitioning		
<b>Structured Streaming and Delta</b>	Streaming Query	Aggregating Streams	Delta Lake		

# Lesson Objectives

By the end of this course, you should be able to:

- 1 Identify core features of Spark and Databricks
- 2 Describe how DataFrames are created and evaluated in Spark
- 3 Apply the DataFrame API to process and analyze data
- 4 Demonstrate how Spark is optimized and executed on a cluster
- 5 Apply Delta Lake and Structured Streaming to process data

Module 1

# Introductions

5



# Welcome!

Let's get to know you



- Name
- Role and team
- Programming experience
- Motivation for attending
- Personal interest

Module 2

# Spark Core





# Spark Core

Databricks Ecosystem

Spark Overview

Spark SQL

Reader & Writer

DataFrame & Column

---

# Databricks Ecosystem

9





# Lakehouse

One simple platform to unify all of  
your data, analytics, and AI workloads

## Customers

**5000+**

across the globe



Original creators of:





Data  
Lake

# Lakehouse

One platform to unify all of  
your data, analytics, and AI  
workloads



Data  
Warehouse



## Data Lake



**DELTA LAKE**

An open approach to bringing  
**data management and**  
governance to data lakes

Better reliability with transactions

48x faster data processing with  
indexing

Data governance at scale with  
fine-grained access control lists



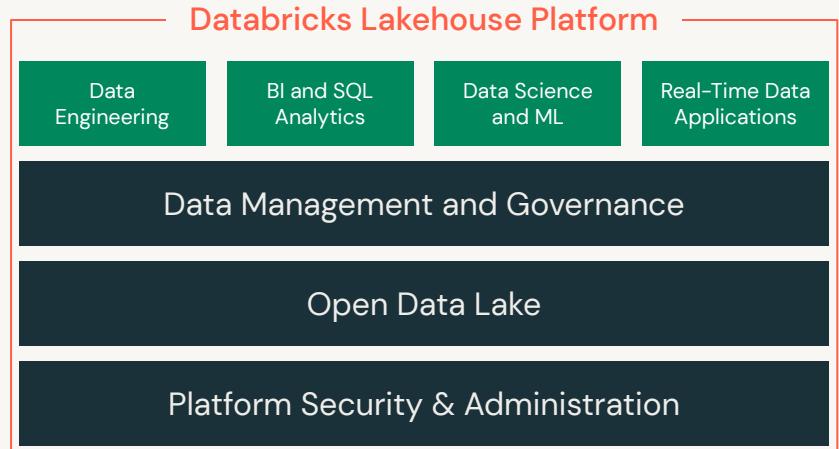
## Data Warehouse

# The Databricks Lakehouse Platform

 Simple

 Open

 Collaborative



Unstructured, semi-structured, structured, and streaming data



# The Databricks Lakehouse Platform



Simple

Unify your data, analytics, and AI on one common platform for all data use cases

## Databricks Lakehouse Platform

Data Engineering

BI and SQL Analytics

Data Science and ML

Real-Time Data Applications

Data Management and Governance

Open Data Lake

Platform Security & Administration



Unstructured, semi-structured, structured, and streaming data



Microsoft Azure



Google Cloud



# The Databricks Lakehouse Platform



Open

Unify your data ecosystem  
with open source standards  
and formats.

Built on the innovation of  
some of the most  
successful open source  
data projects in the world

**30 Million+**  
Monthly downloads



mlflow™



re'dash

# The Databricks Lakehouse Platform

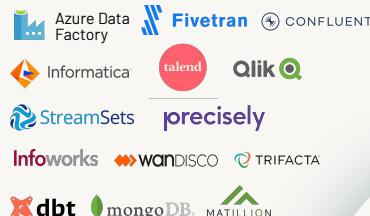


Unify your data ecosystem  
with open source standards  
and formats.

450+

Partners across the  
data landscape

## Visual ETL & Data Ingestion



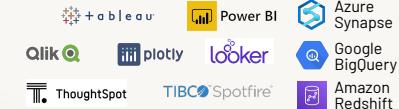
## Data Providers



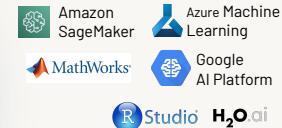
## Top Consulting & SI Partners



## Business Intelligence



## Machine Learning



## Centralized Governance



databricks  
Lakehouse Platform

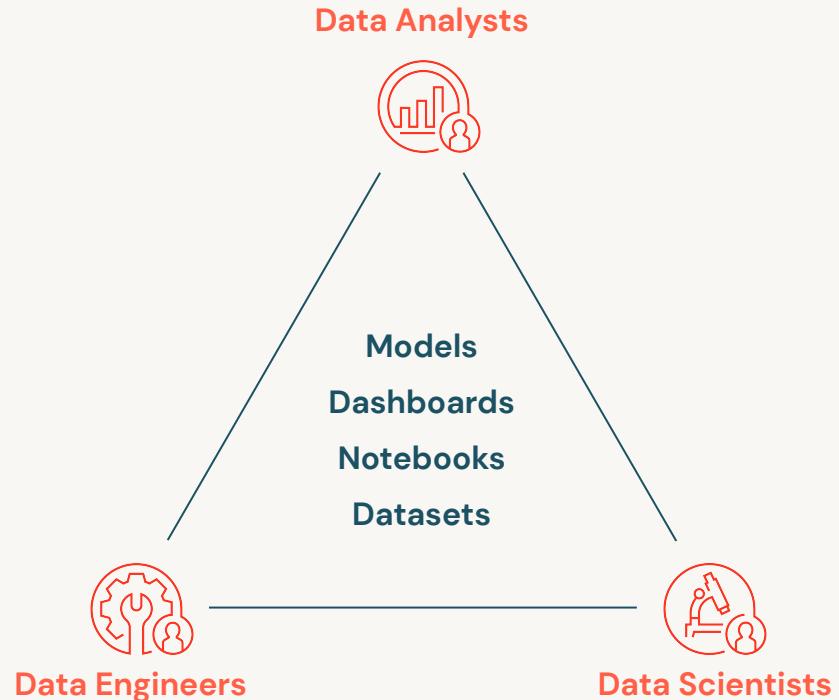
Microsoft Azure    aws  
Google Cloud

# The Databricks Lakehouse Platform



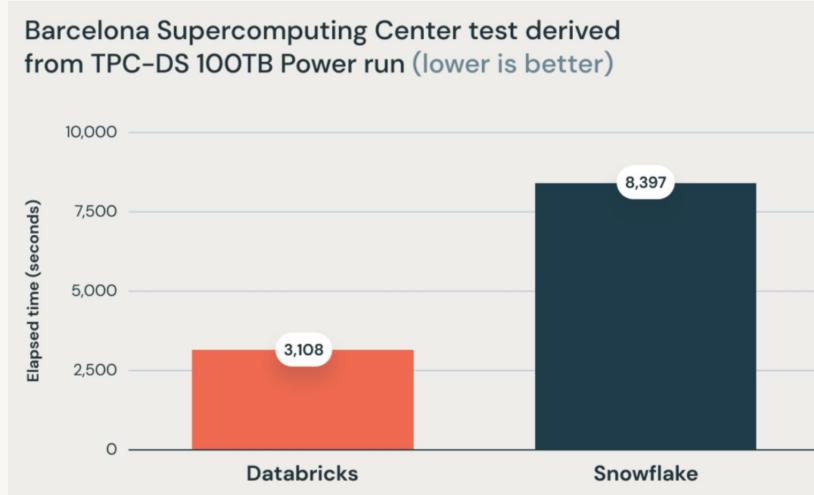
## Collaborative

Unify your data teams to collaborate across the entire data and AI workflow



# TPC-DS

Databricks SQL set **official data warehousing performance record** – outperformed the previous record by 2.2x.



# Spark Overview



- De-facto standard unified analytics engine for big data processing
- Largest open-source project in data processing
- Technology created by the founders of Databricks

# Spark Benefits



Fast



Easy to Use



Unified



# Spark API

Spark SQL +  
DataFrames

Streaming

MLlib

Spark Core API

R

SQL

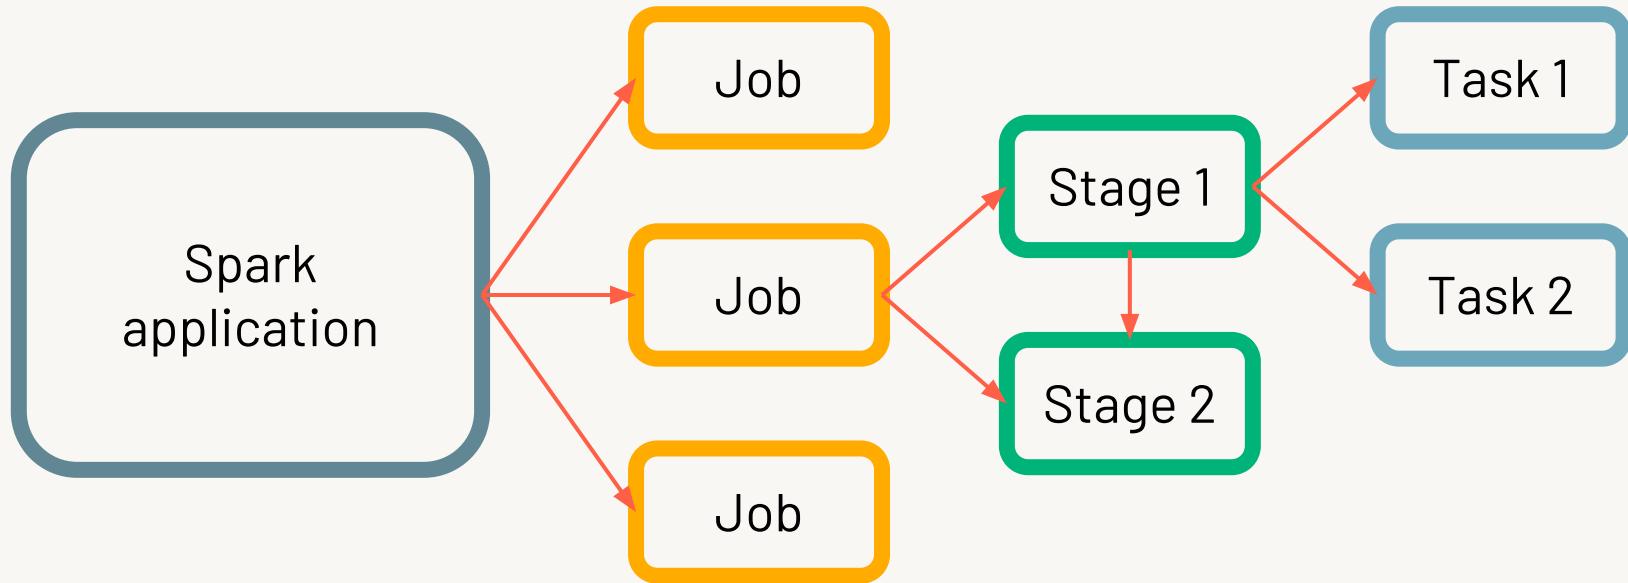
Python

Scala

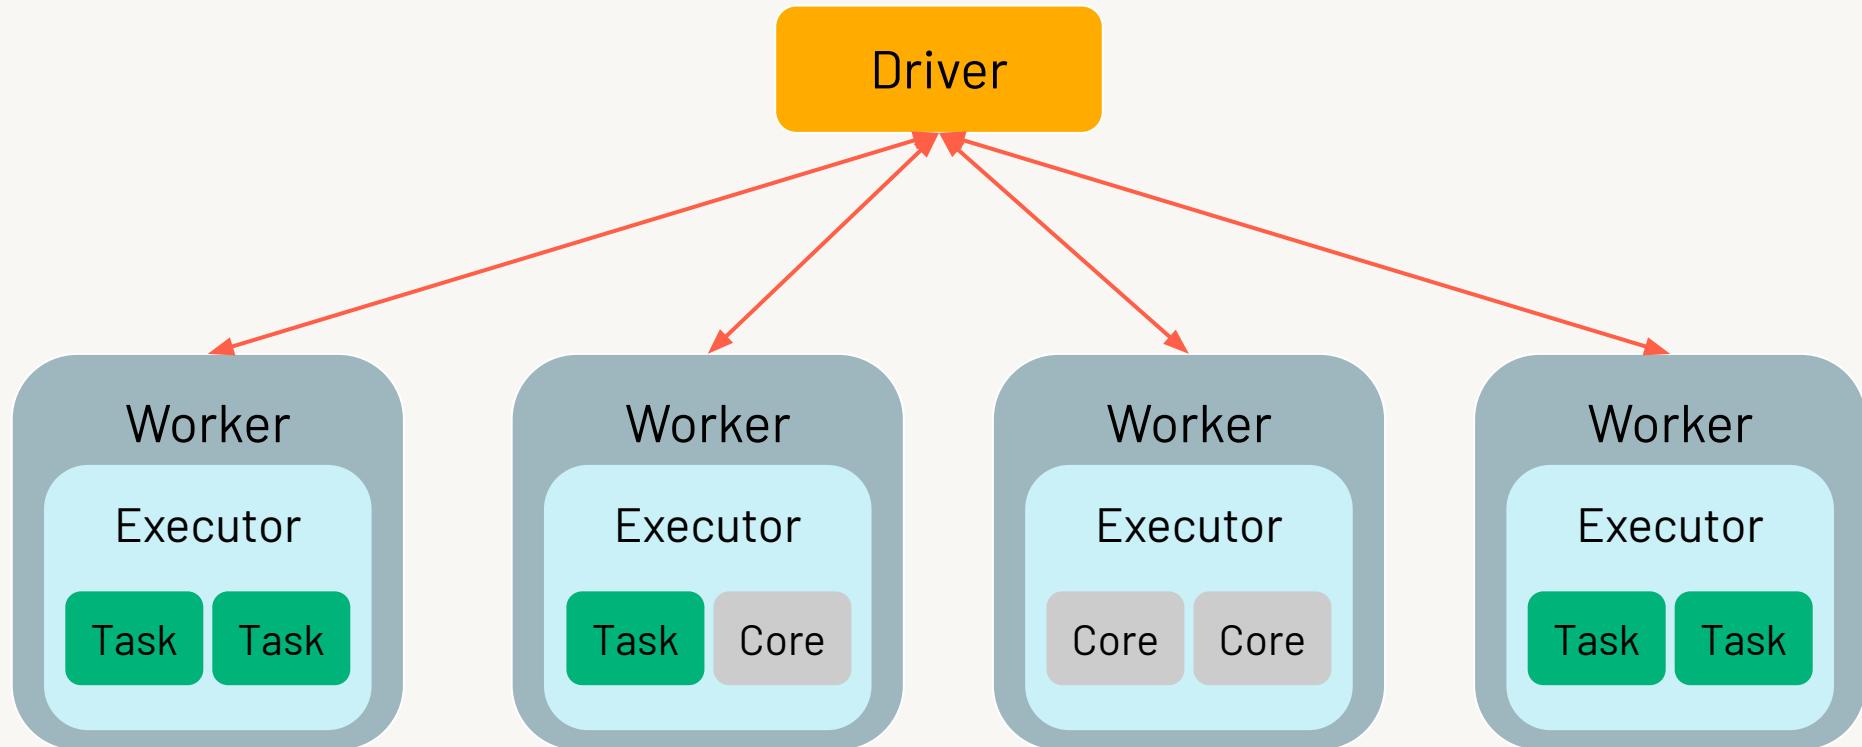
Java



# Spark Execution



# Spark Cluster



# CSV

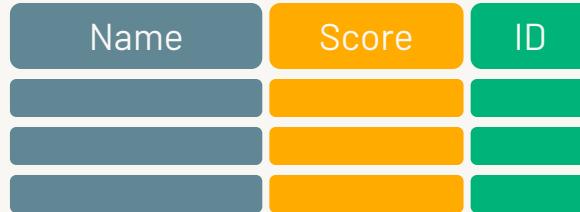
## Comma Separated Values

```
item_id, name, price, qty
M_PREM_Q, Premium Queen Mattress, 1795, 35
M_STAN_F, Standard Full Mattress, 945, 24
M_PREM_F, Premium Full Mattress, 1695, 45
M_PREM_T, Premium Twin Mattress, 1095, 18
```



# Parquet

A columnar storage format



Row-Oriented data on disk



Column-Oriented data on disk

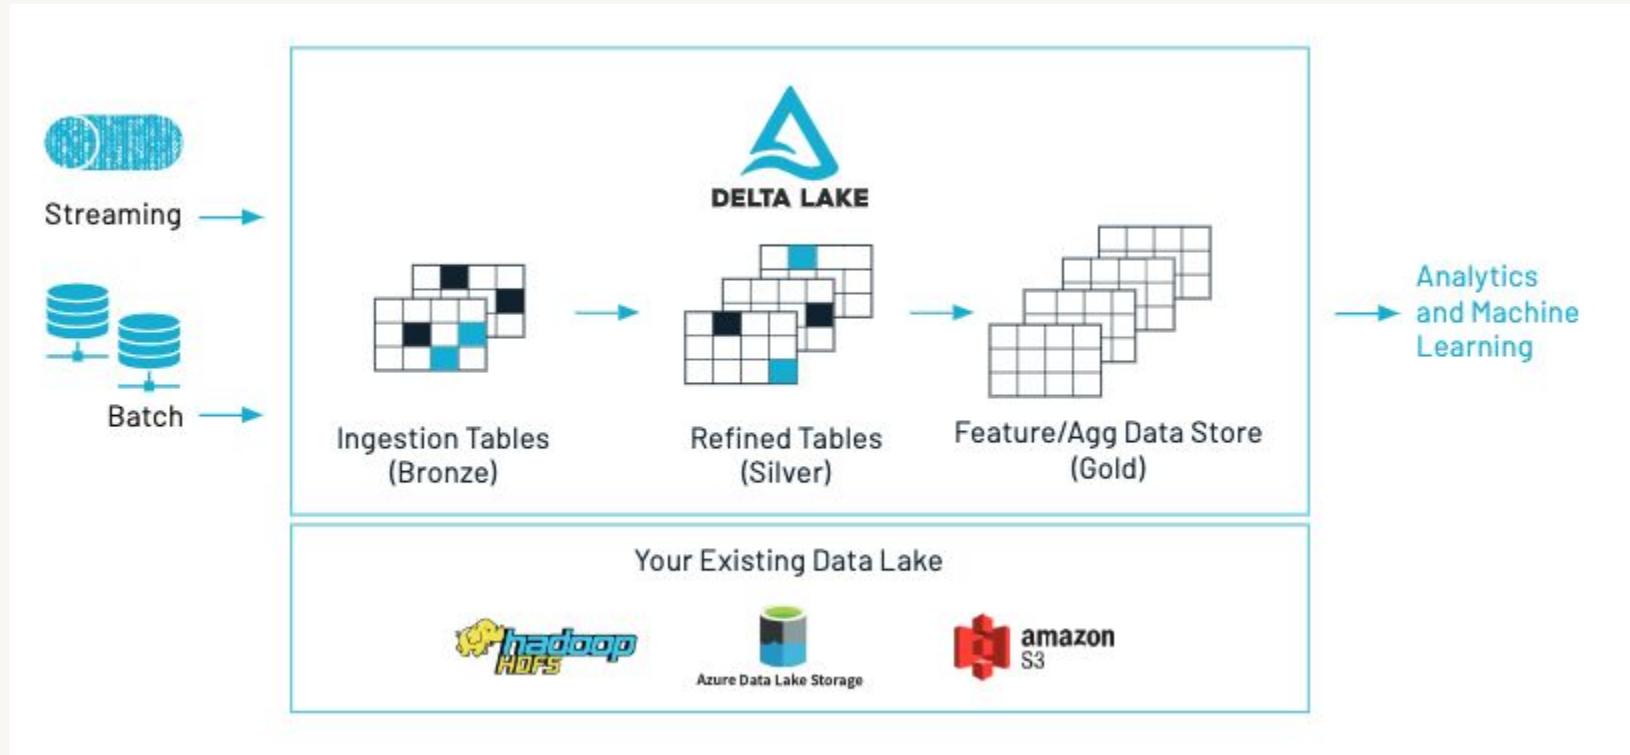


# Delta Lake

Technology designed to be used with Apache Spark to build robust data lakes



# Open-source Storage Layer



# Delta Lake's Key Features

- ACID transactions
- Time travel (data versioning)
- Schema enforcement and evolution
- Audit history
- Parquet format
- Compatible with Apache Spark API



Module 3

# Functions



# Functions

Aggregation

Datetimes

Complex Types

Additional Functions

User-Defined Functions

---

Module 4

# Performance





# Performance

Spark Architecture

Query Optimization

Partitioning



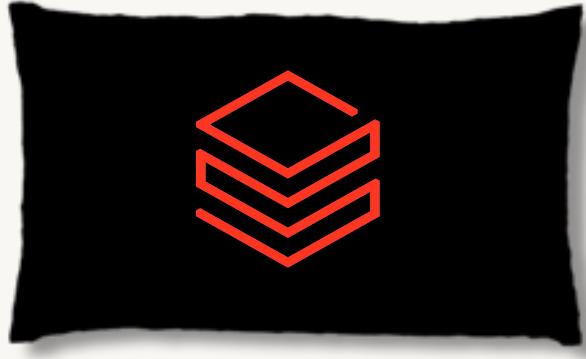
# Spark Architecture

Spark Cluster

Spark Execution

34





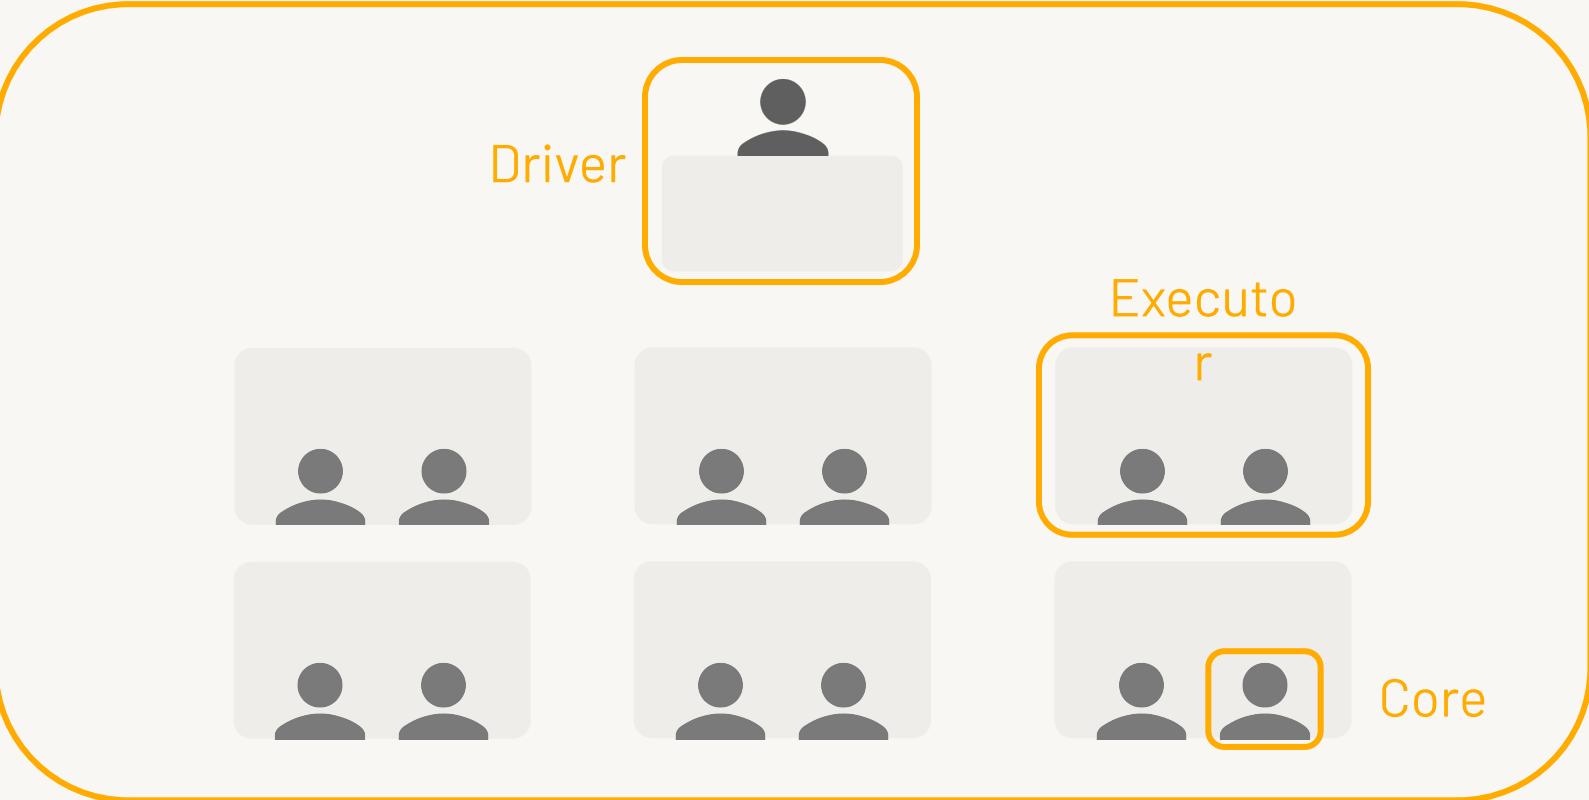
**Scenario 1: Filter out brown pieces from candy bags**

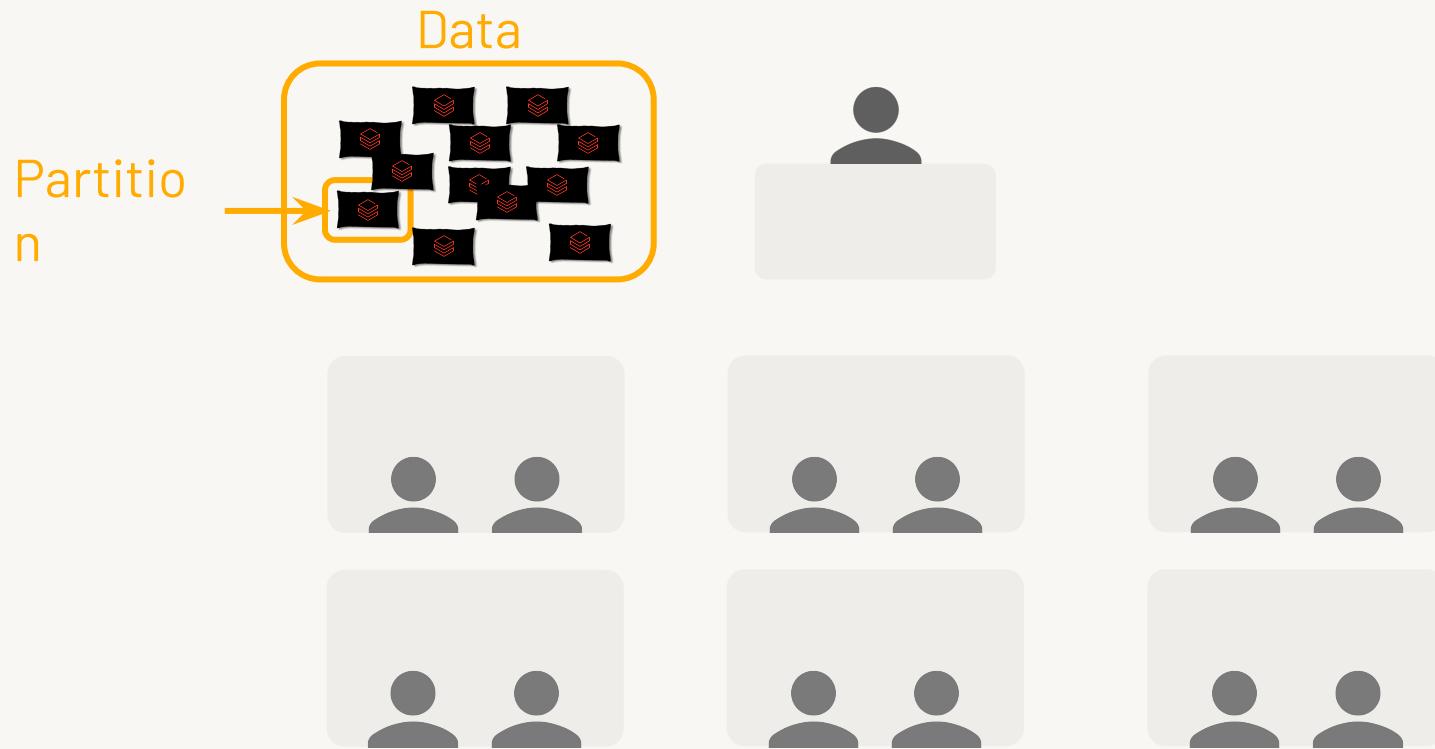
Cluster

Driver

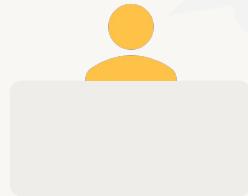
Executor

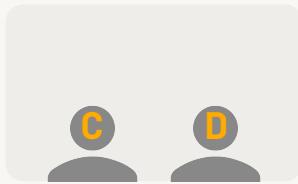
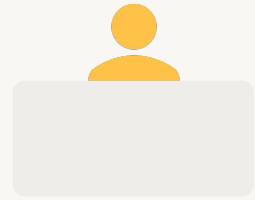
Core

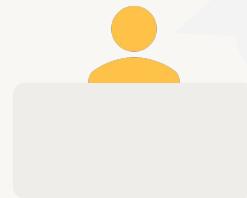




We need filter out brown pieces  
from these candy bags







Student A get bag # 1  
Student B get bag # 2  
Student C get bag # 3

...



A  
B



C  
D



E  
F



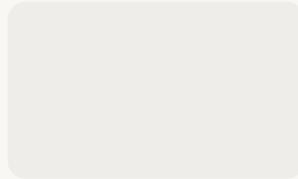
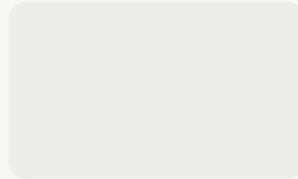
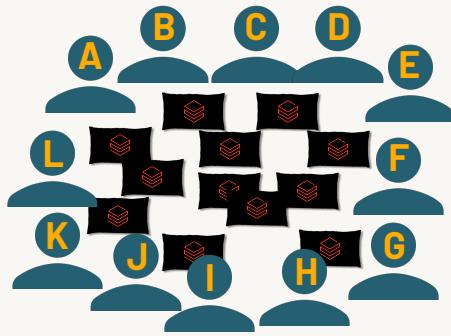
G  
H

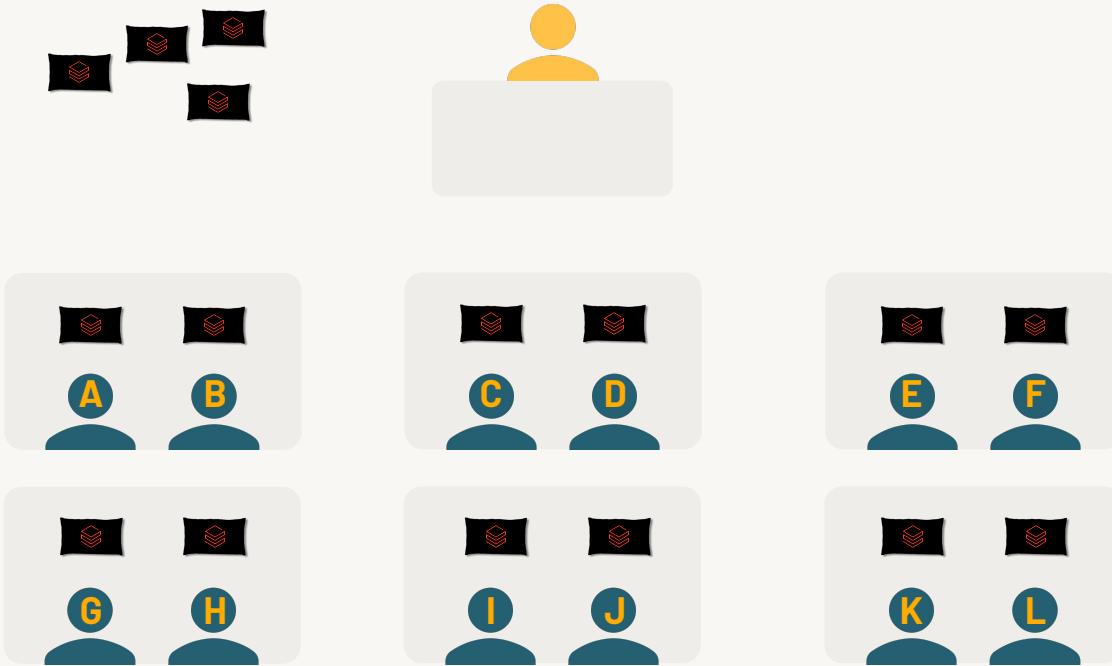


I  
J

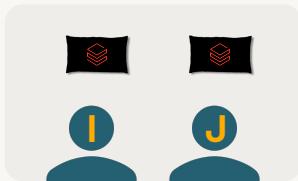
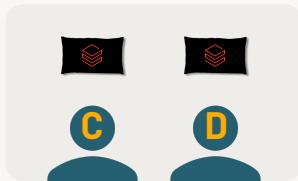
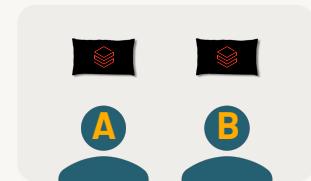
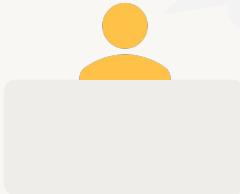


K  
L

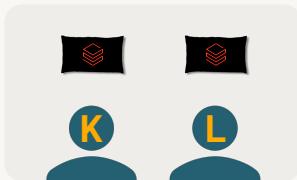
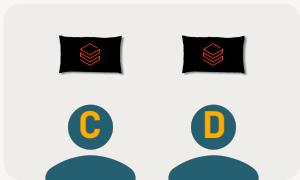
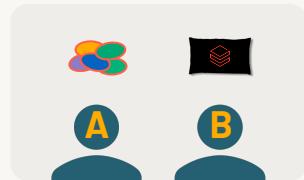
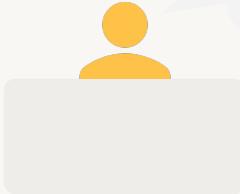


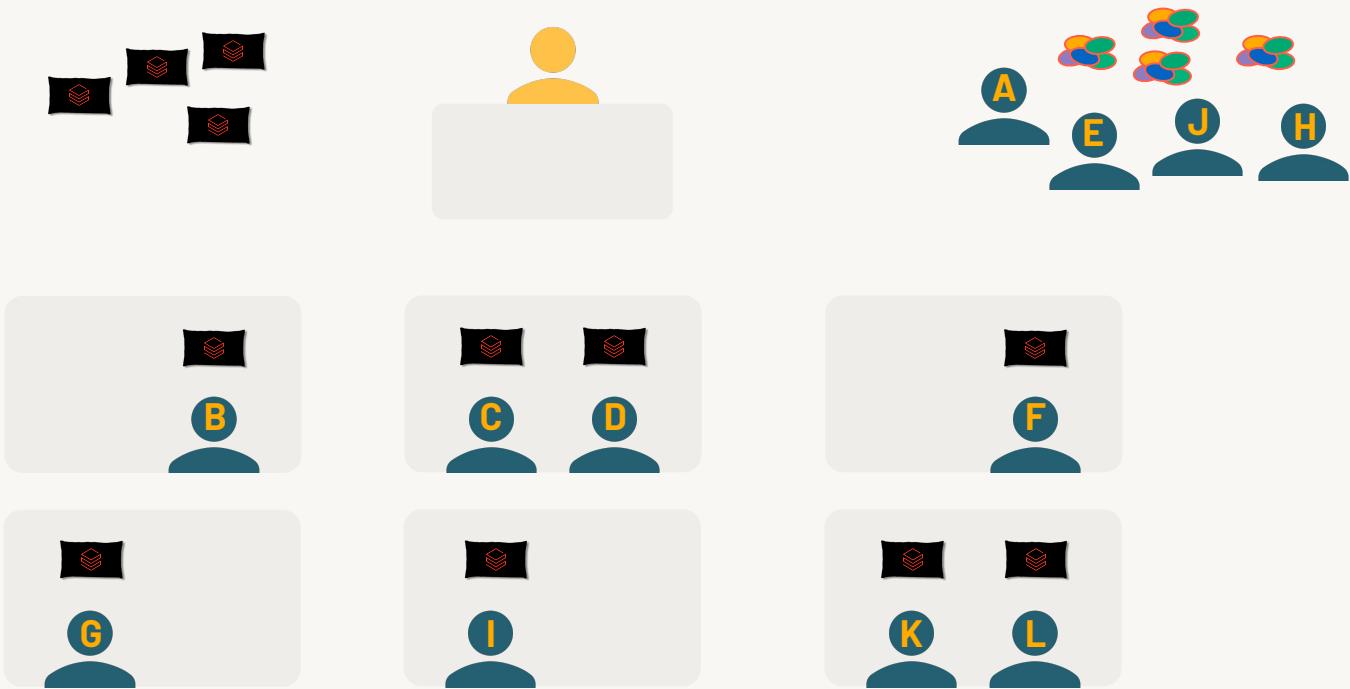


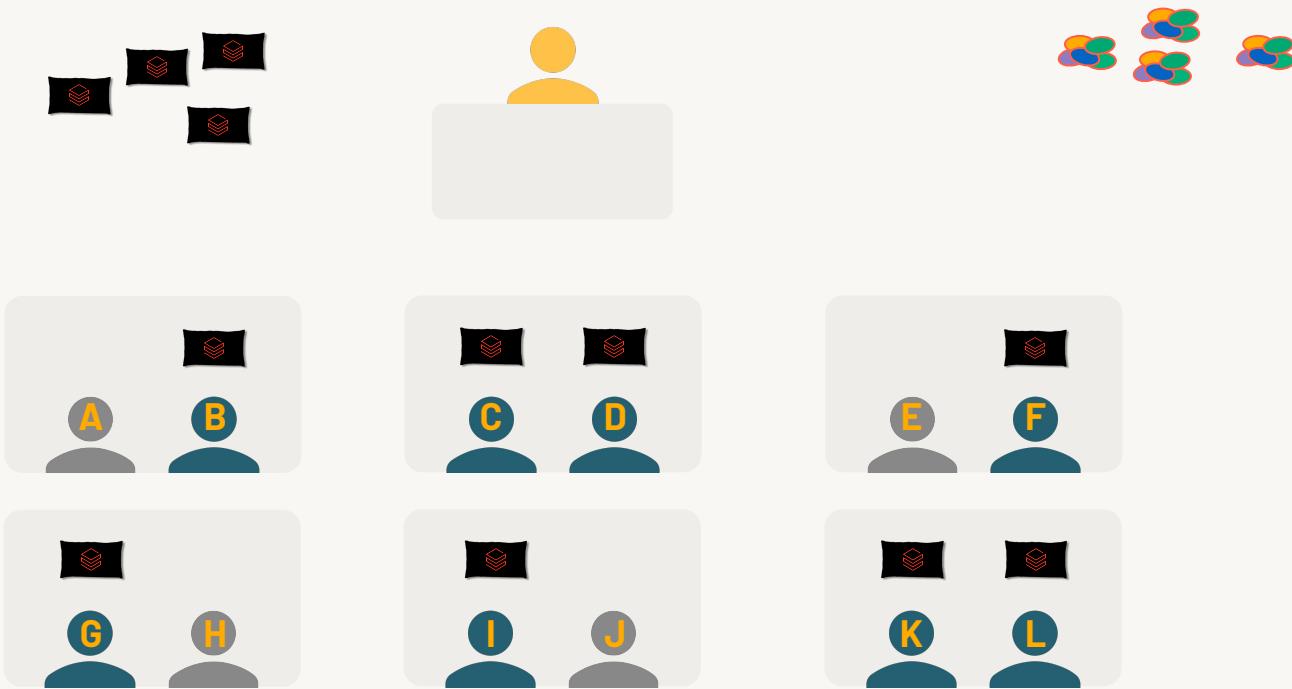
Eliminate the brown candy pieces  
and pile the rest in the corner



Eliminate the brown candy pieces  
and pile the rest in the corner

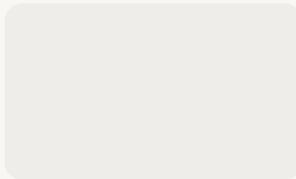
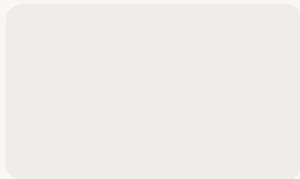
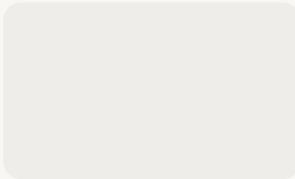
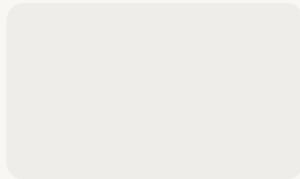
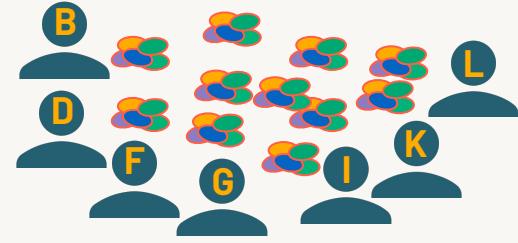
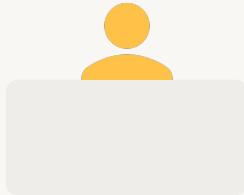
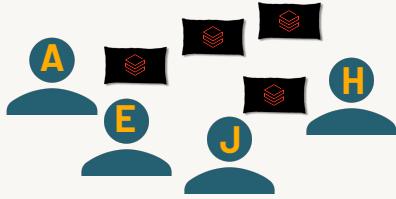


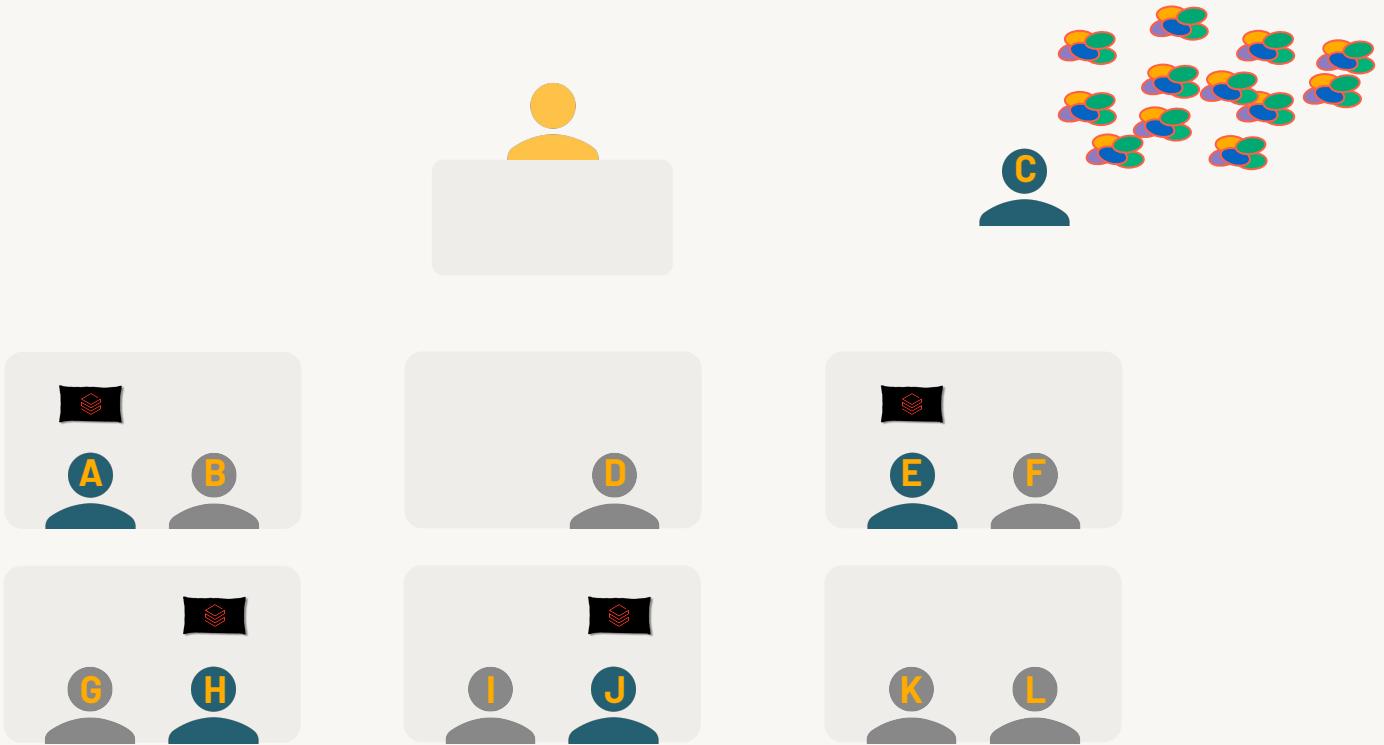


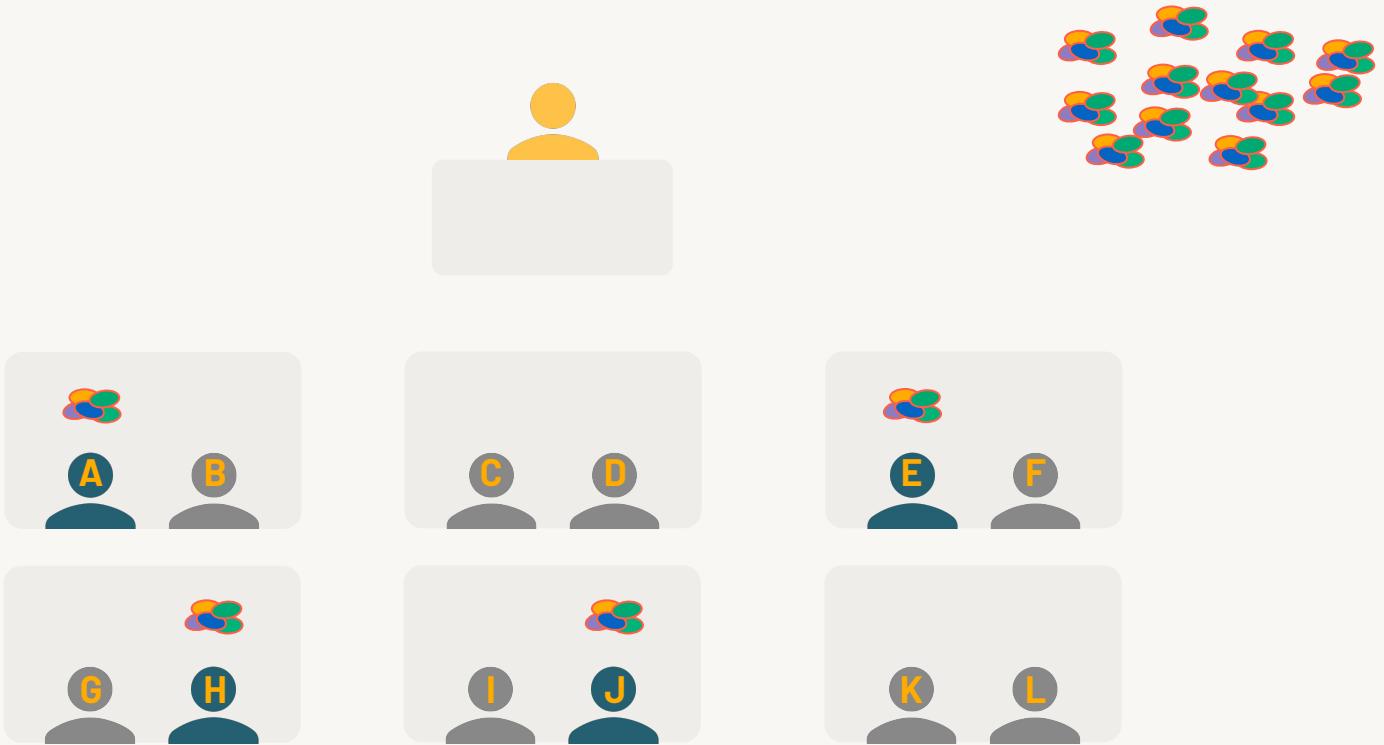


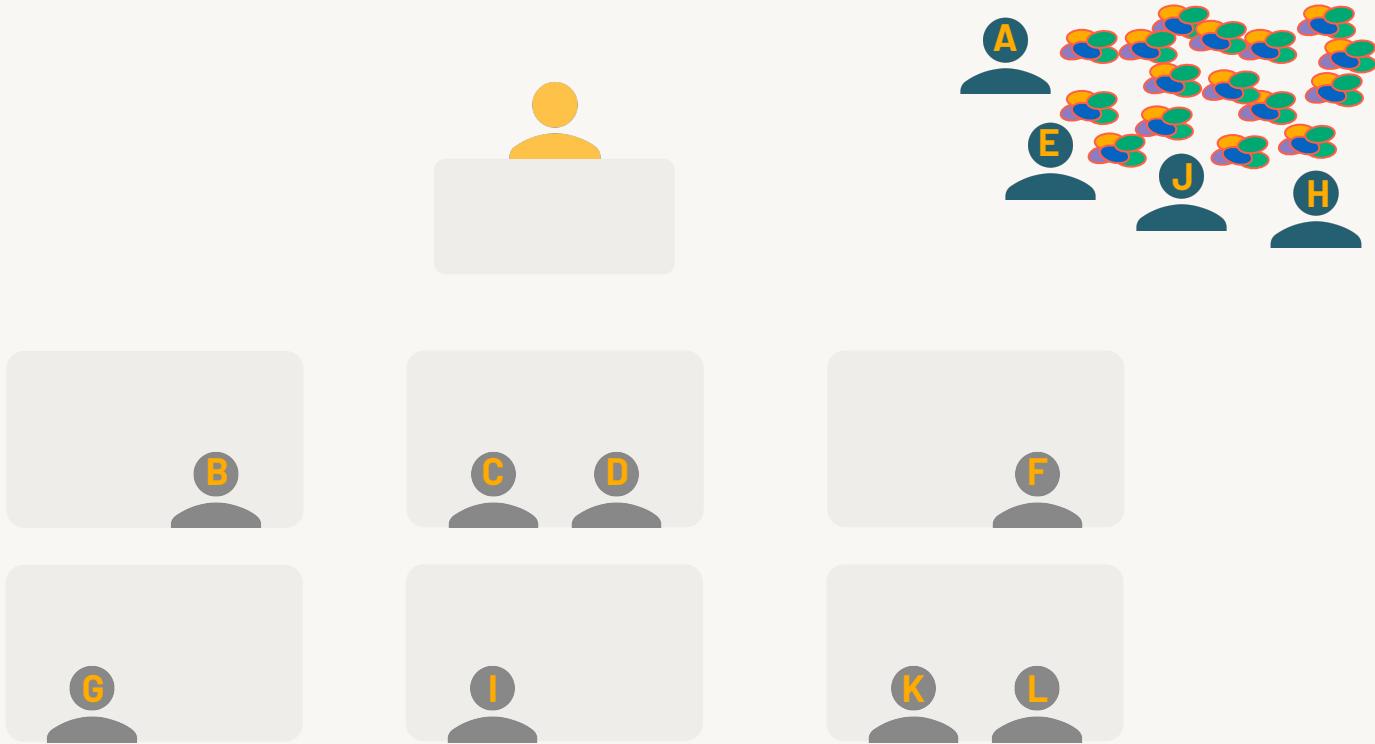
Students A, E, H, J,  
get bags 13, 14, 15, 16

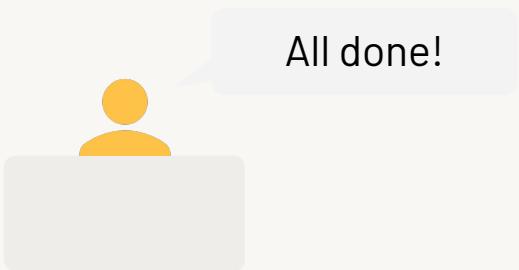




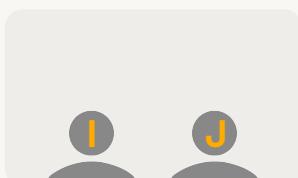
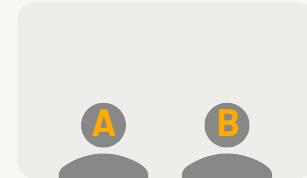
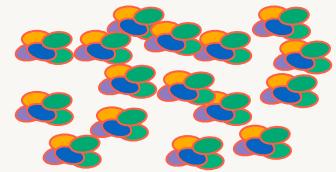


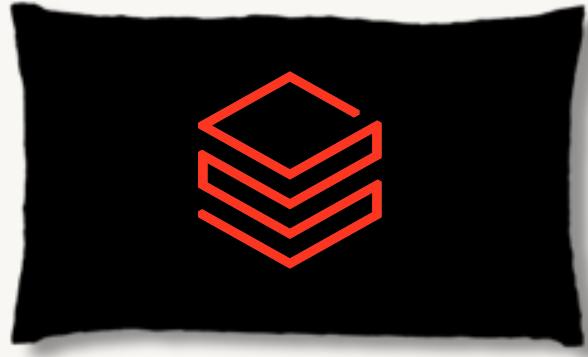






All done!





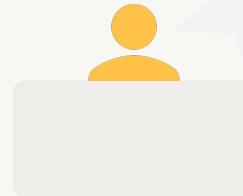
## Scenario 2: Count total pieces in candy bags



# Stage 1: Local Count



# Stage 1: Local Count



We need to count the total pieces  
in these candy bags



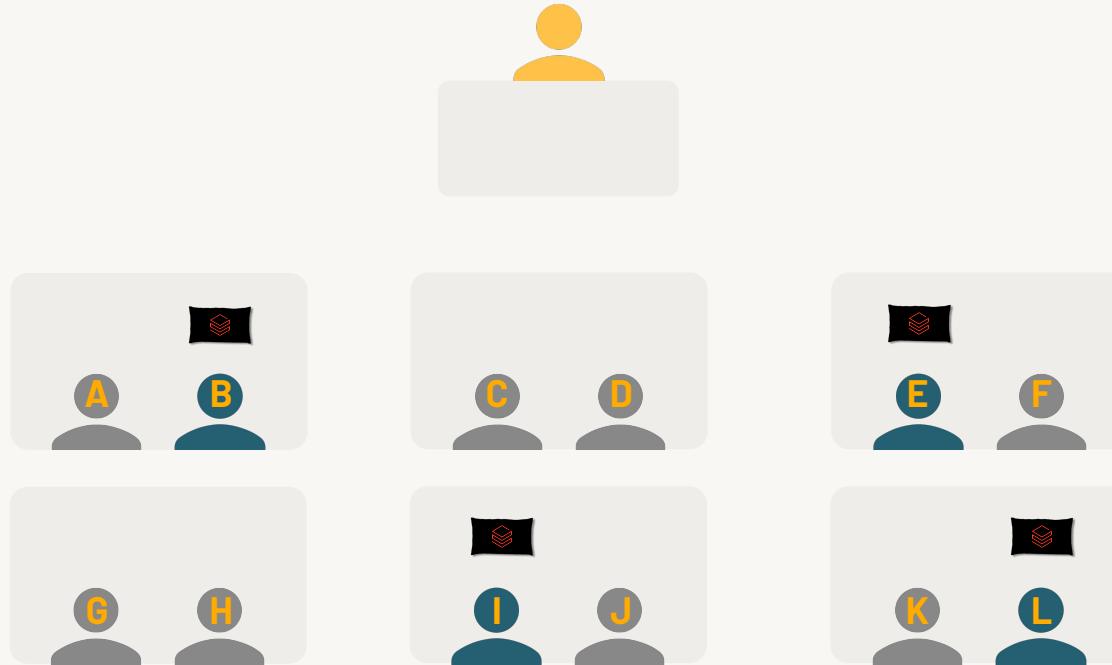
# Stage 1: Local Count



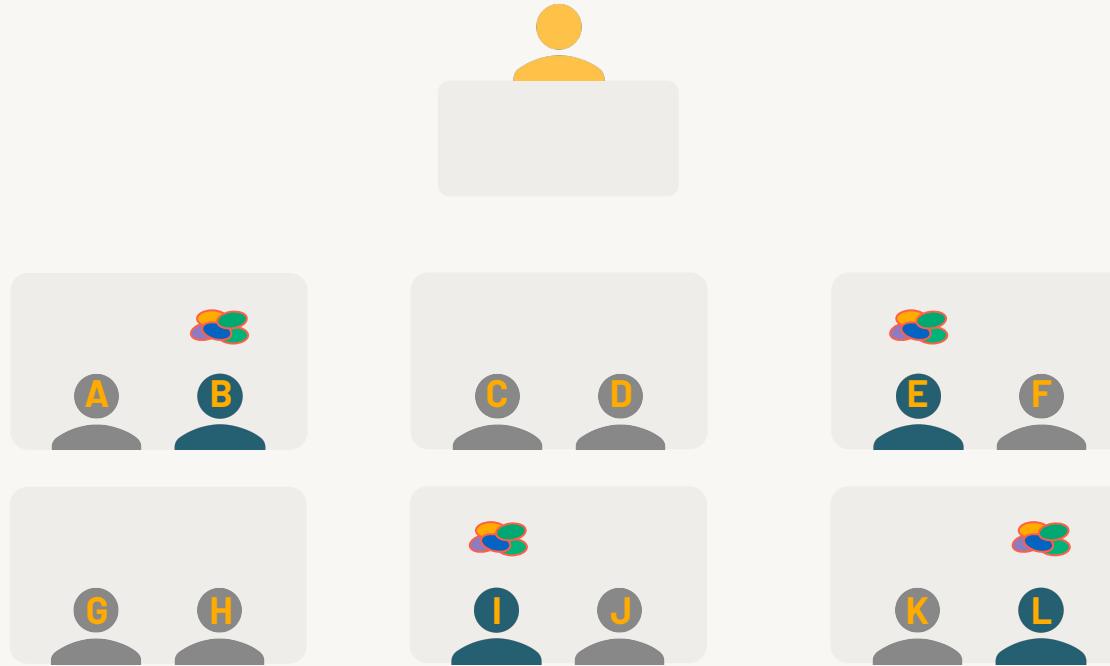
Students B, E, I, L,  
get these four bags



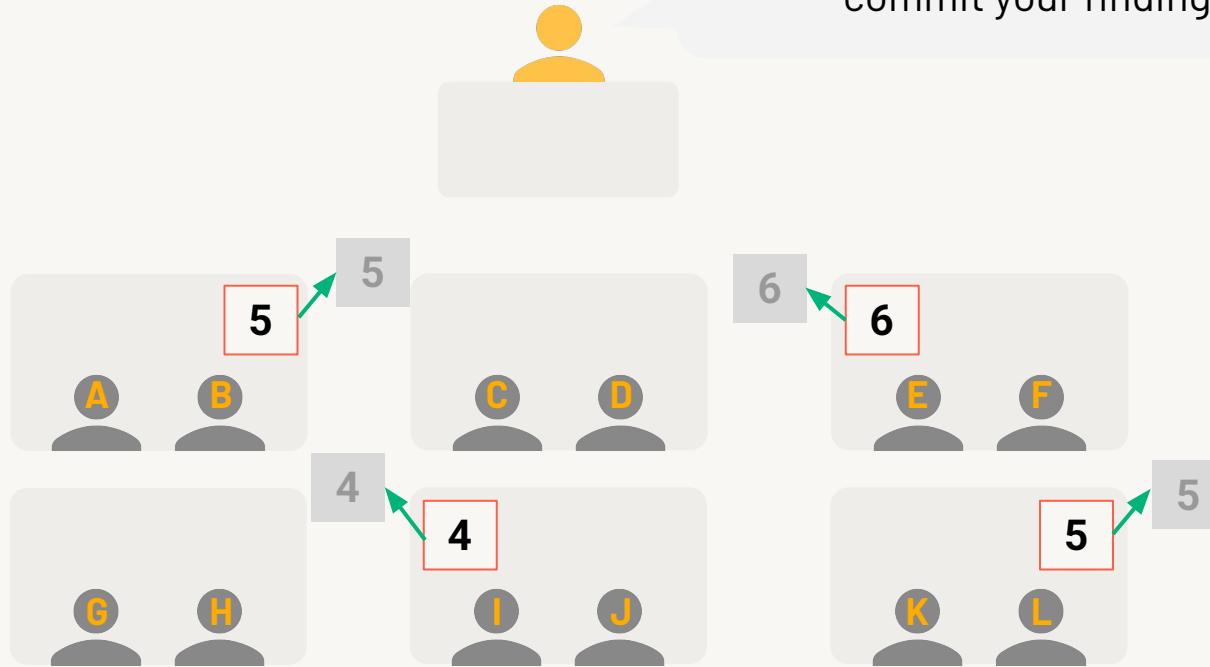
# Stage 1: Local Count



# Stage 1: Local Count



# Stage 1: Local Count



# Stage 1: Local Count

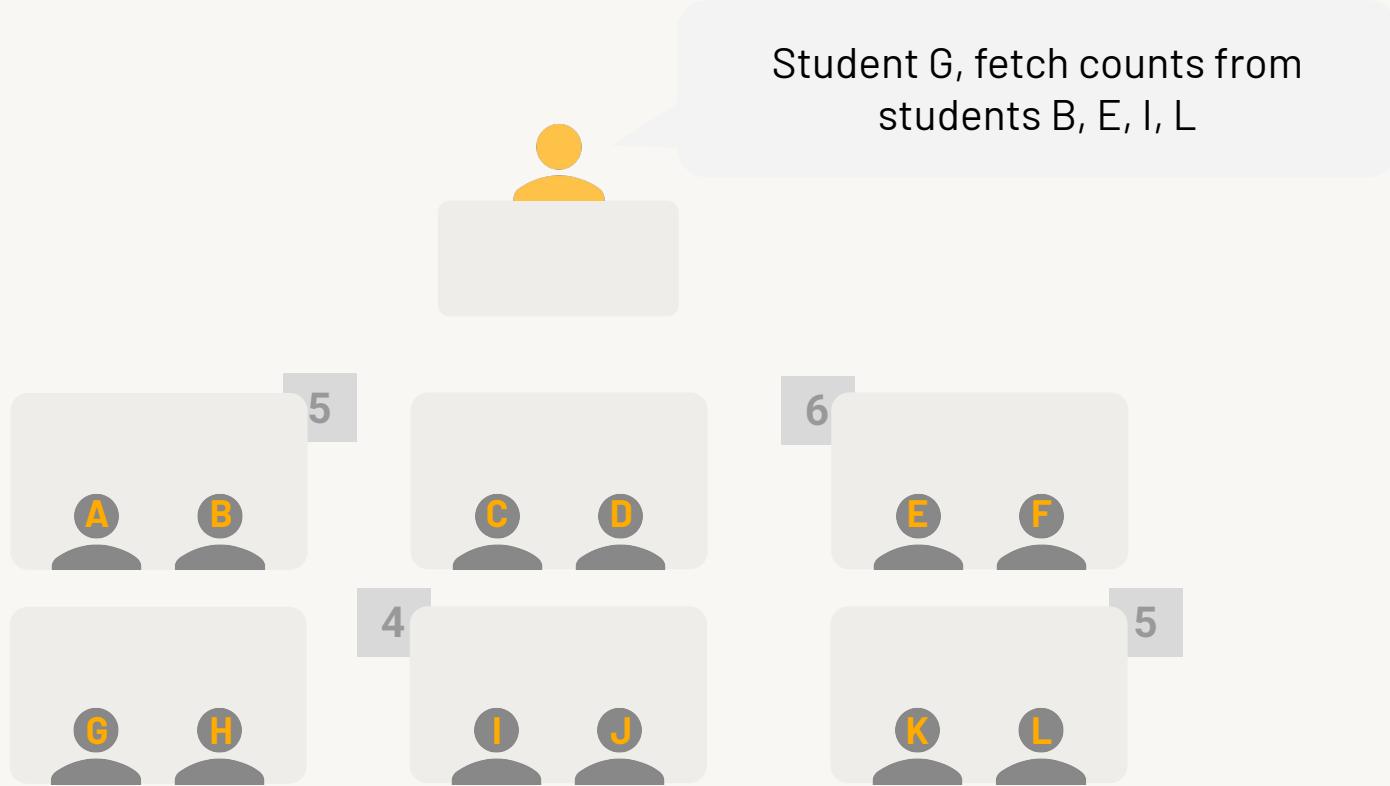
Stage 1 is complete!



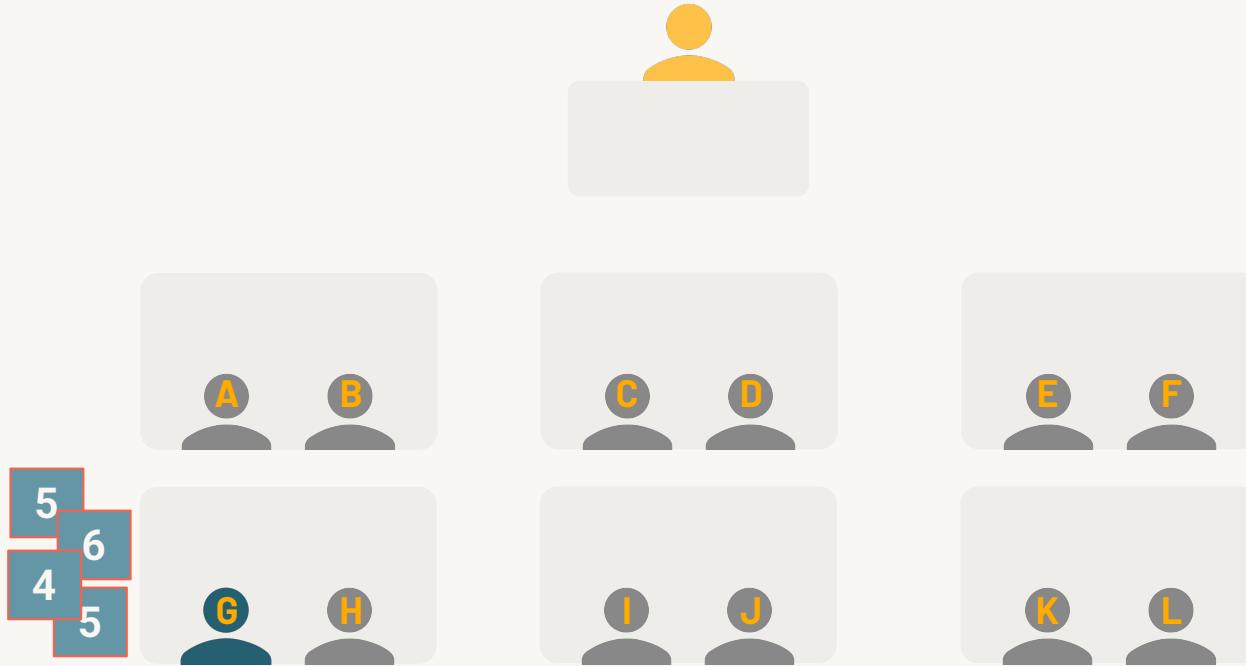
# Stage 2: Global Count



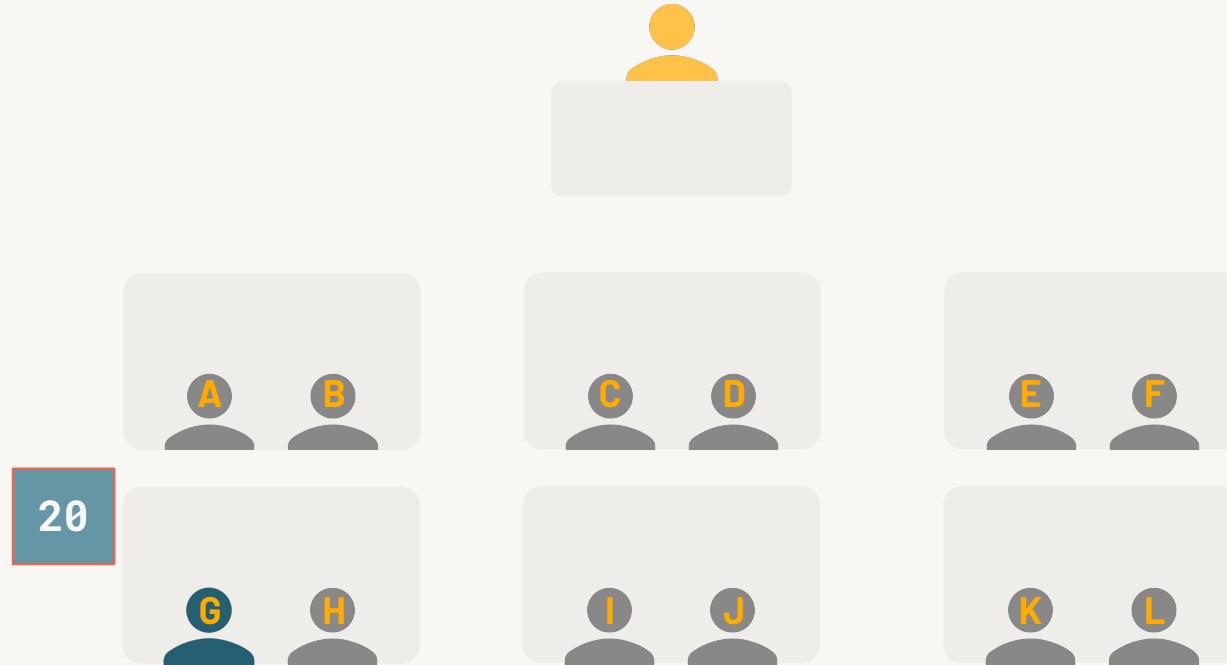
## Stage 2: Global Count



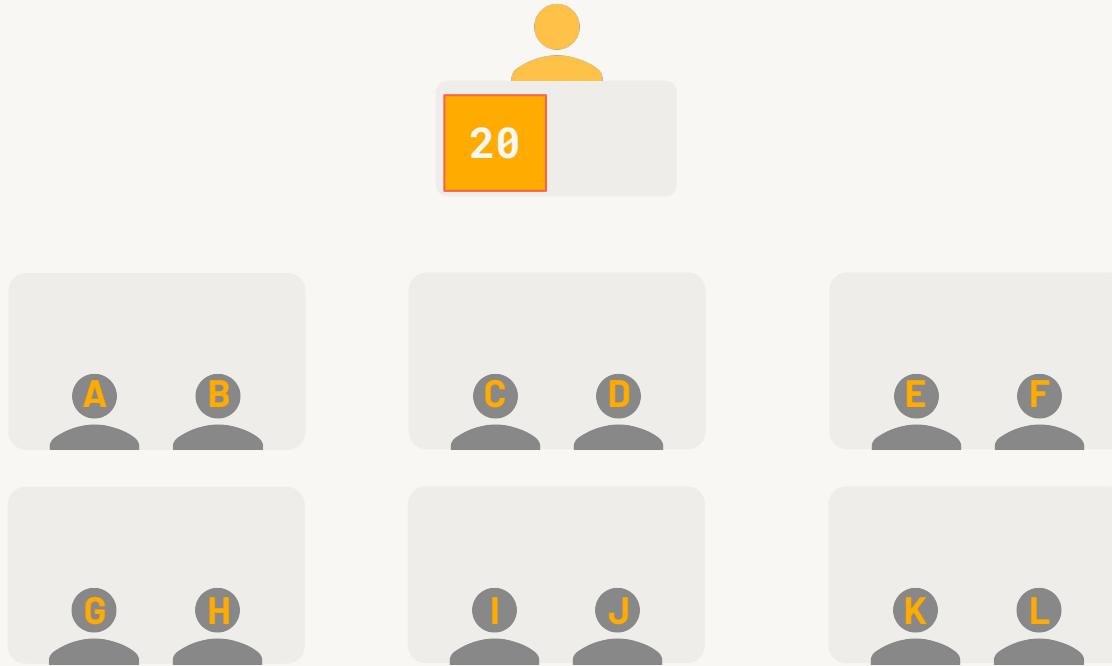
## Stage 2: Global Count



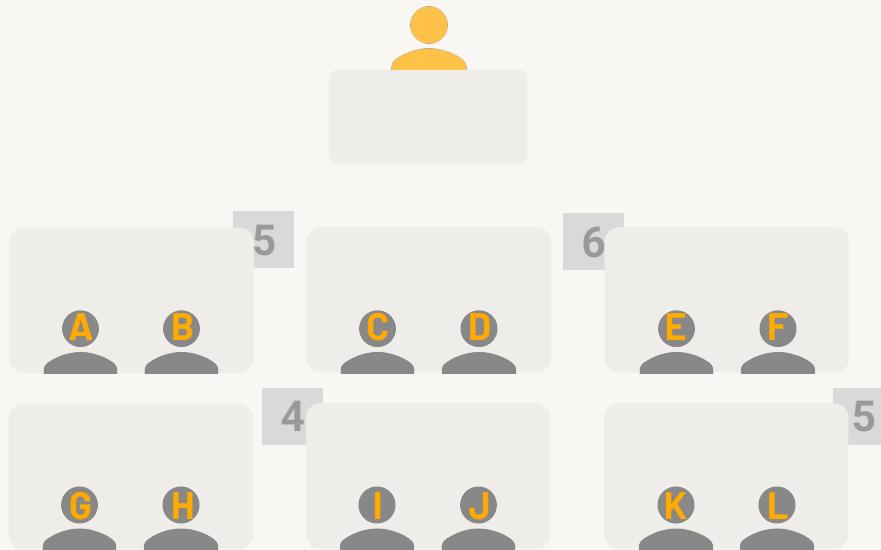
## Stage 2: Global Count



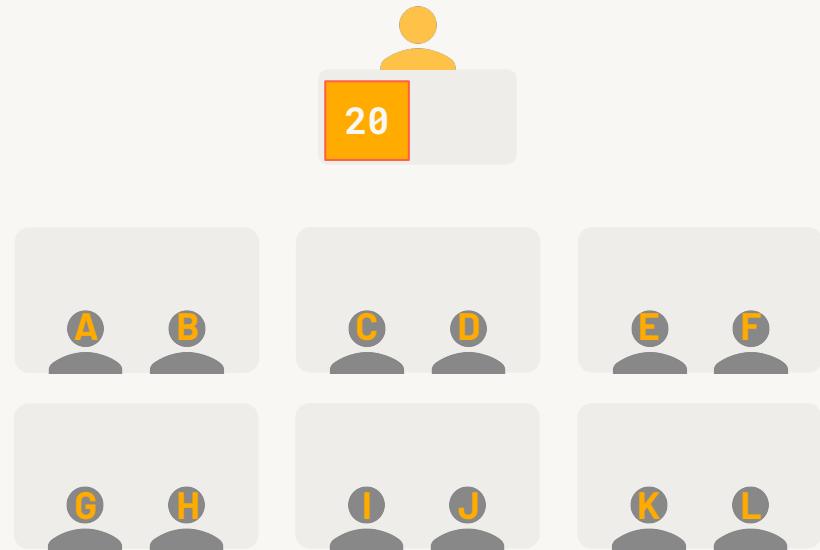
## Stage 2: Global Count



## Stage 1: Local Count



## Stage 2: Global Count



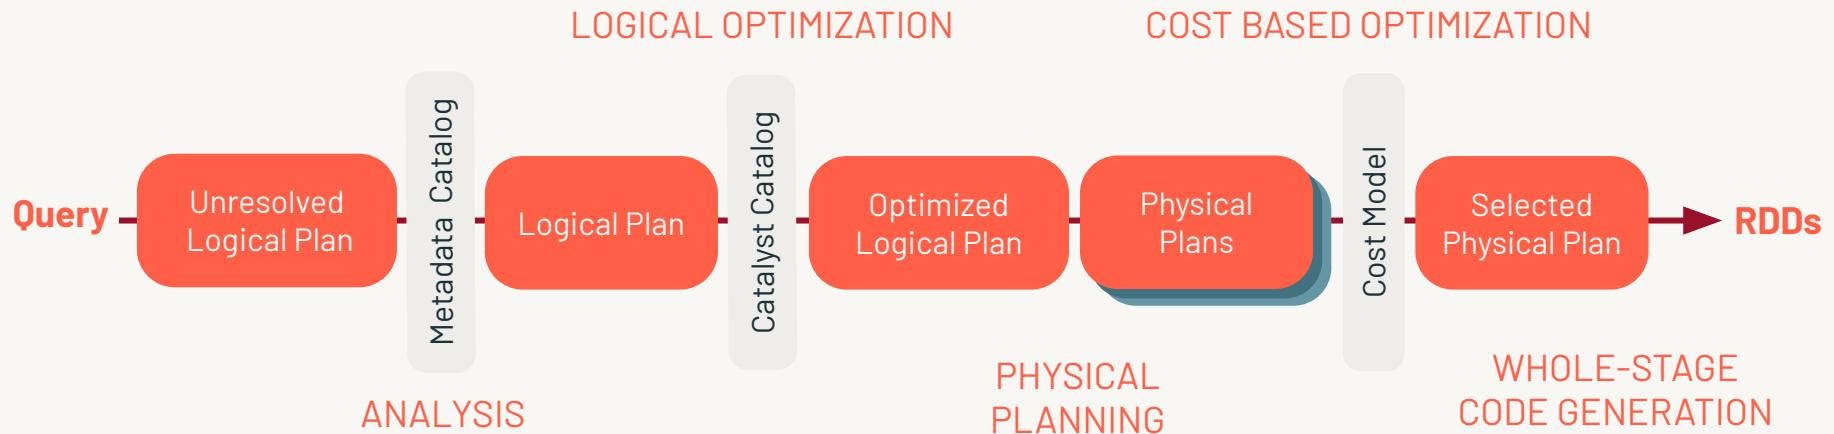
# Query Optimization

Catalyst Optimizer  
Adaptive Query Execution

67

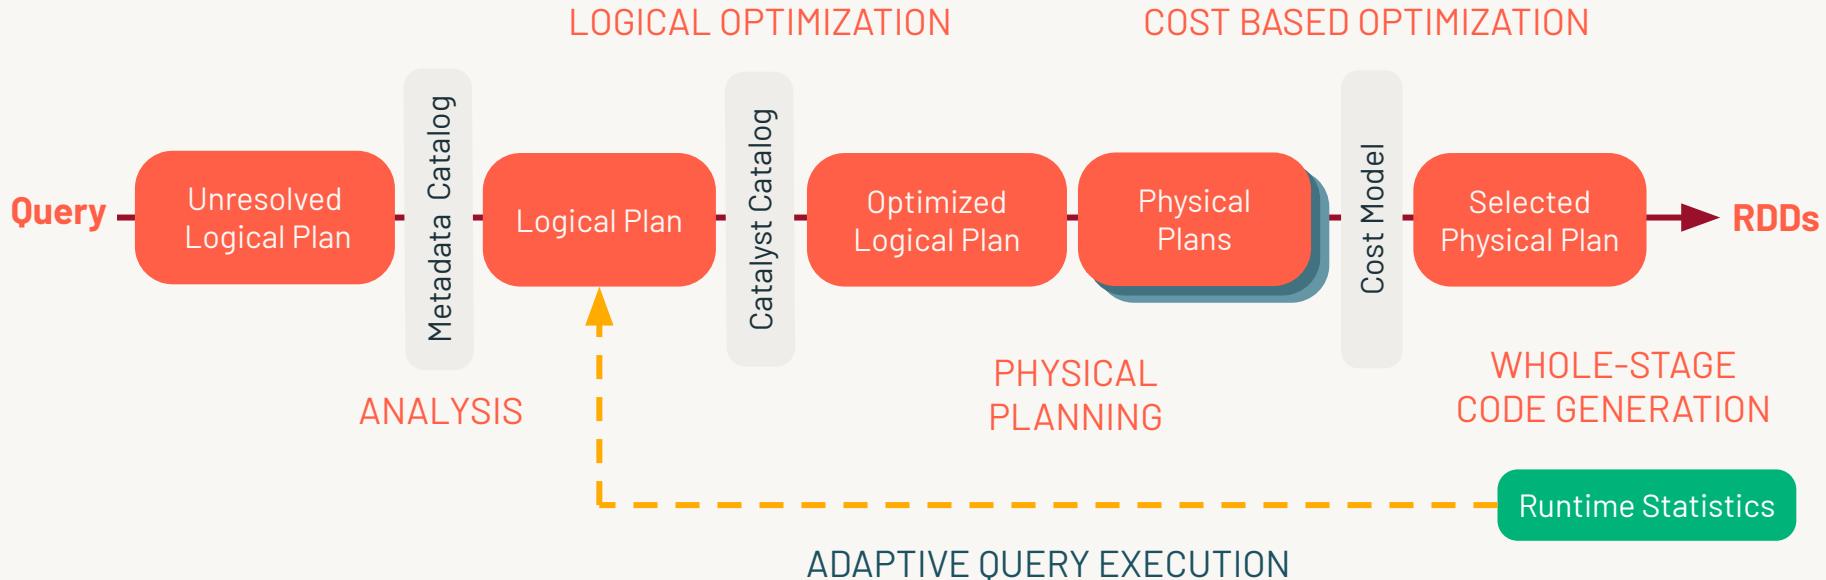


# Query Optimization



# Query Optimization with AQE

New in Spark 3.0, **enabled** by default as of Spark 3.2



Module 5

# Structured Streaming





# Structured Streaming

Streaming Query  
Stream Aggregations

---

# Streaming Query

Advantages

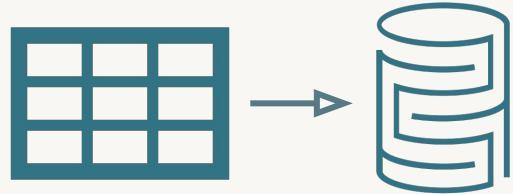
Use Cases

Sources

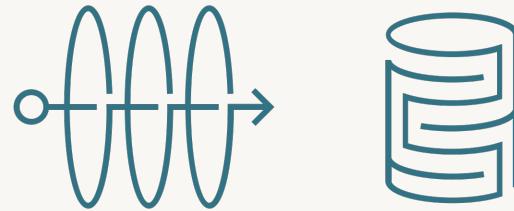
72



# Batch Processing



# Stream Processing



# Advantages of Stream Processing



Lower latency



Efficient Updates

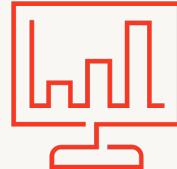


Automatic bookkeeping on new data

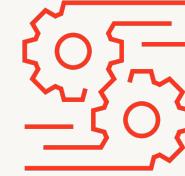
# Stream Processing Use Cases



Notifications



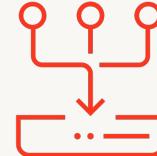
Real-time  
reporting



Incremental ETL



Update data to  
serve in  
real-time

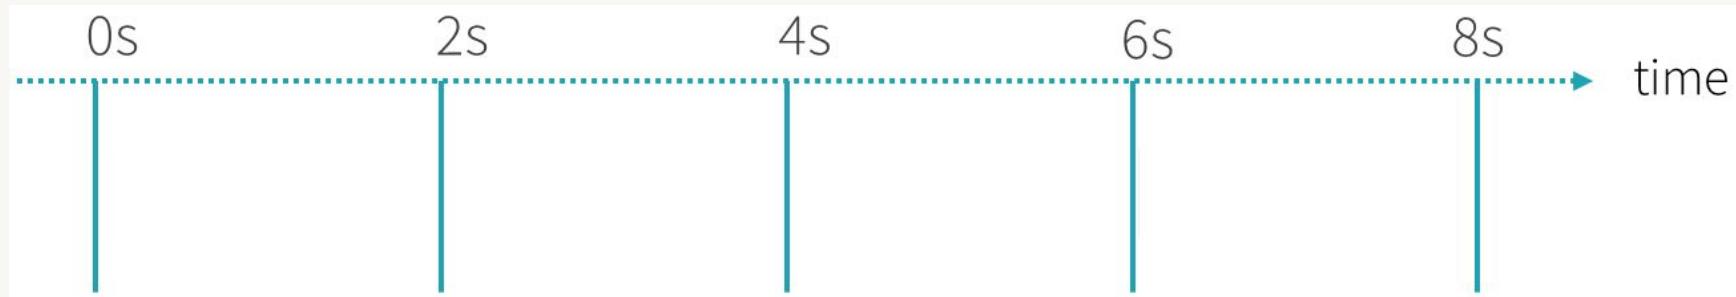


Real-time  
decision making

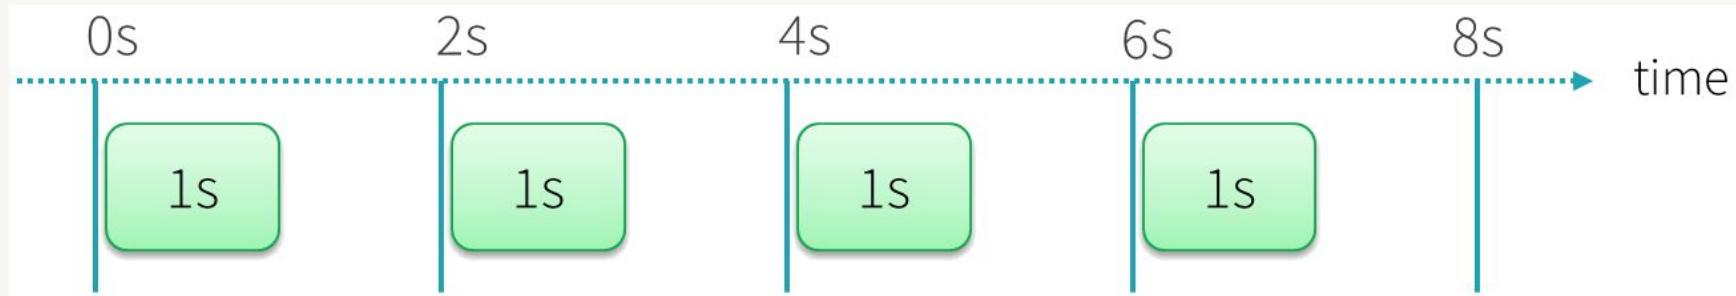


Online ML

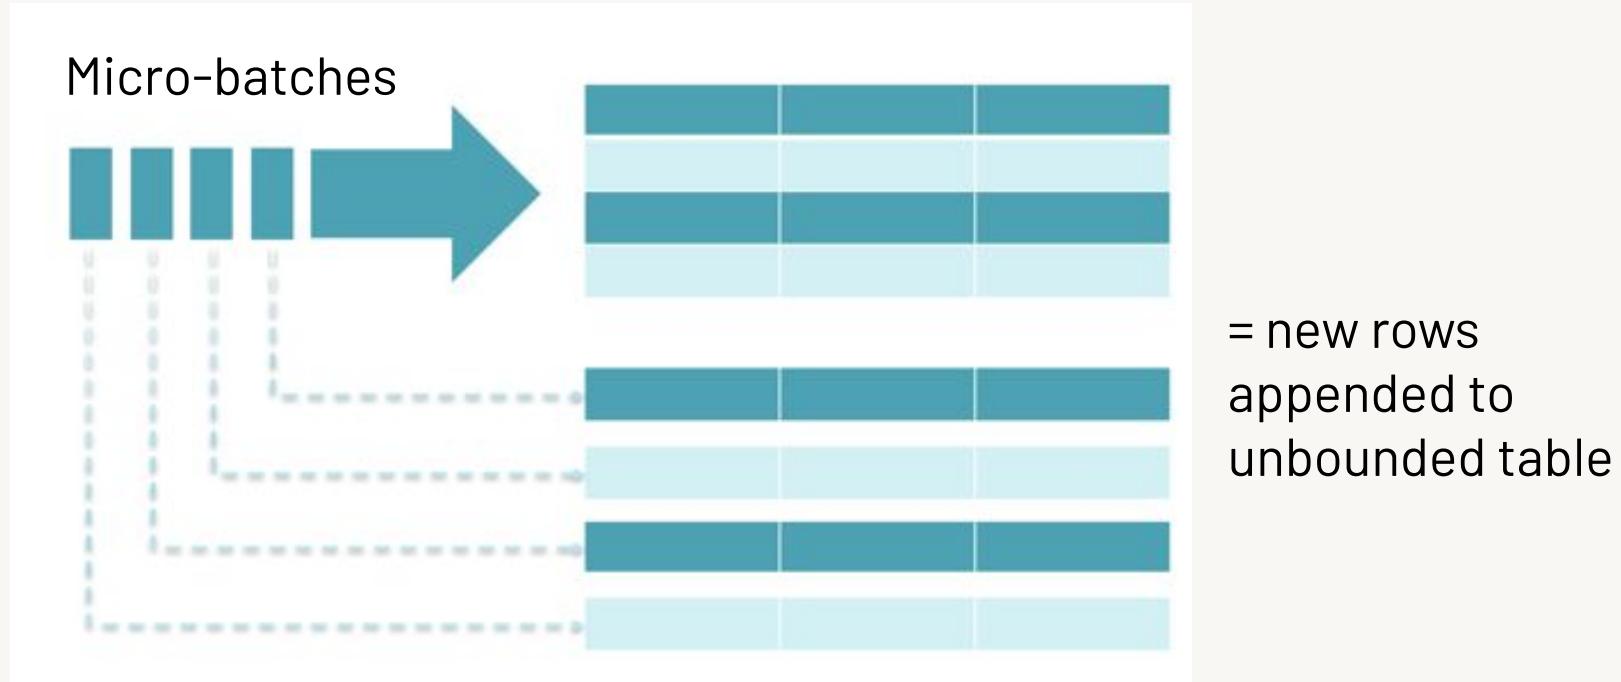
# Micro-Batch Processing



# Micro-Batch Processing



# Structured Streaming

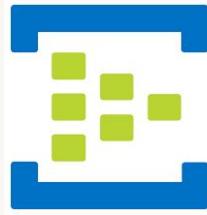


# Input Sources

Kafka



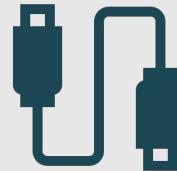
Event Hubs



Files



Sockets



Generator



FOR TESTING

# Sinks

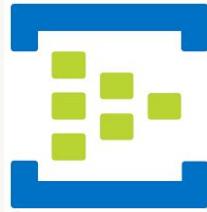
Kafka



Files



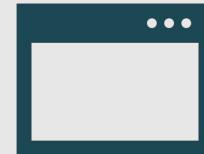
Event Hubs



Foreach



Console



Memory



FOR DEBUGGING

# Output Modes

## APPEND

Add new records  
only

## UPDATE

Update changed  
records in place

## COMPLETE

Rewrite full output



# Trigger Types

<b>Default</b>	Process each micro-batch as soon as the previous one has been processed
<b>Fixed interval</b>	Micro-batch processing kicked off at the user-specified interval
<b>One-time</b>	Process all of the available data as a single micro-batch and then automatically stop the query
<b>Continuous Processing</b>	Long-running tasks that continuously read, process, and write data as soon events are available <i>*Experimental</i> See <a href="#">Structured Streaming Programming Guide</a>



# End-to-end fault tolerance

Guaranteed in Structured Streaming by

Checkpointing and write-ahead logs

Idempotent sinks

Replayable data sources



# Stream Aggregations

Aggregations  
Windows  
Watermarking

84



# Real-time Aggregations



Errors in IoT data by device type



Anomalous behavior in server log files by country



Behavior analysis on messages by hashtags

# Time-Based Windows

## Tumbling Windows

**No window overlap**

Any given event gets aggregated into **only one** window group

e.g. 1:00–2:00 am, 2:00–3:00 am, 3:00–4:00 am, ...

## Sliding Windows

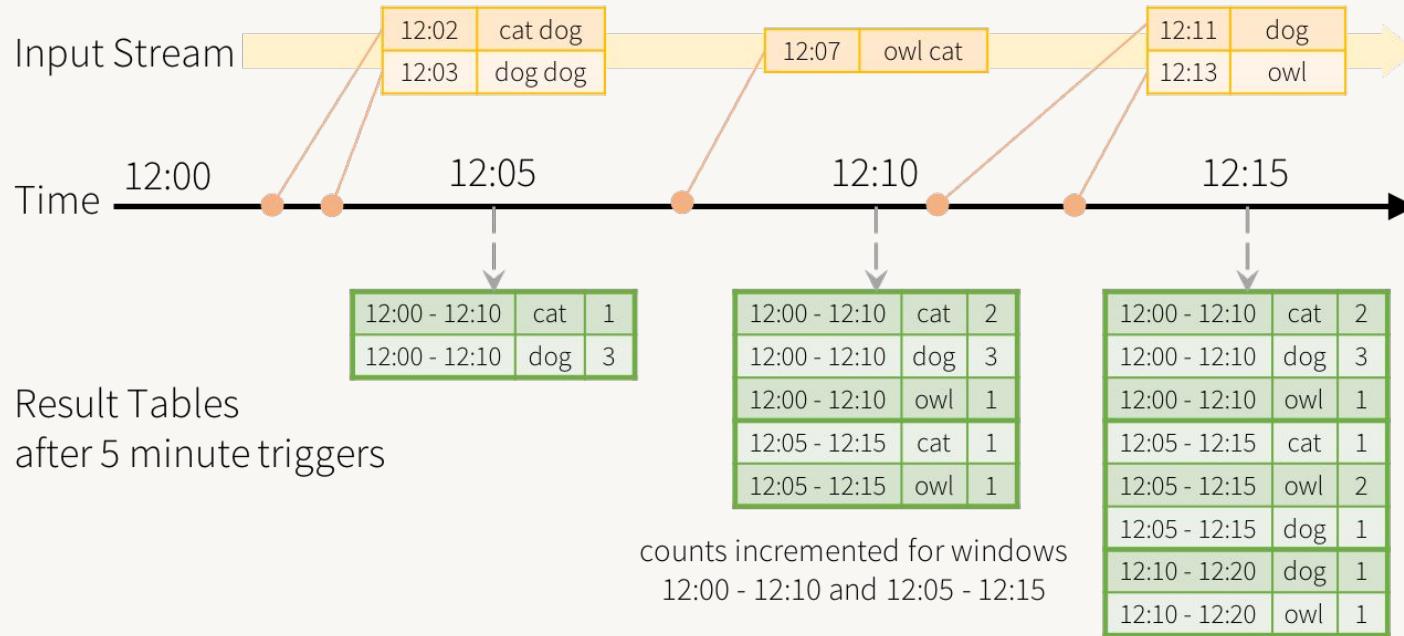
**Windows overlap**

Any given event gets aggregated into **multiple** window groups

e.g. 1:00–2:00 am, 1:30–2:30 am, 2:00–3:00 am, ...



# Sliding Windows Example



# Windowing

```
(streamingDF
    .groupBy(col("device"),
              window(col("time"), "1 hour"))
    .count())
```

[Cancel](#)

▼ (1) Spark Jobs

▼ Job 4 [View](#) (3 stages)

Stage 12:	<div style="width: 100%;">1/1 (0 running)</div>	<a href="#">i</a>
Stage 13:	<div style="width: 39.5%;">79/200 (4 running)</div>	<a href="#">i</a>
Stage 14:	<div style="width: 0%;">0/1 (0 running)</div>	<a href="#">i</a>

Why are we seeing 200 tasks for this stage?



# Control the Shuffle Repartitioning

```
spark.conf.set("spark.sql.shuffle.partitions", spark.sparkContext.defaultParallelism)
```

The screenshot shows the Databricks UI interface. At the top left, there is a 'Cancel' button and a 'More' menu icon. Below that, a tree view shows '(1) Spark Jobs' expanded, revealing 'Job 42'. Under 'Job 42', there are three stages listed: Stage 126, Stage 127, and Stage 128. Each stage has a progress bar and a status message. Stage 126 shows '1/1 (0 running)'. Stage 127 shows '4/4 (0 running)'. Stage 128 shows '1/1 (0 running)'. To the right of each stage entry is an information icon (a blue circle with a white 'i').

Stage	Status	Info
Stage 126	1/1 (0 running)	i
Stage 127	4/4 (0 running)	i
Stage 128	1/1 (0 running)	i



# Event-Time Processing

## EVENT-TIME DATA

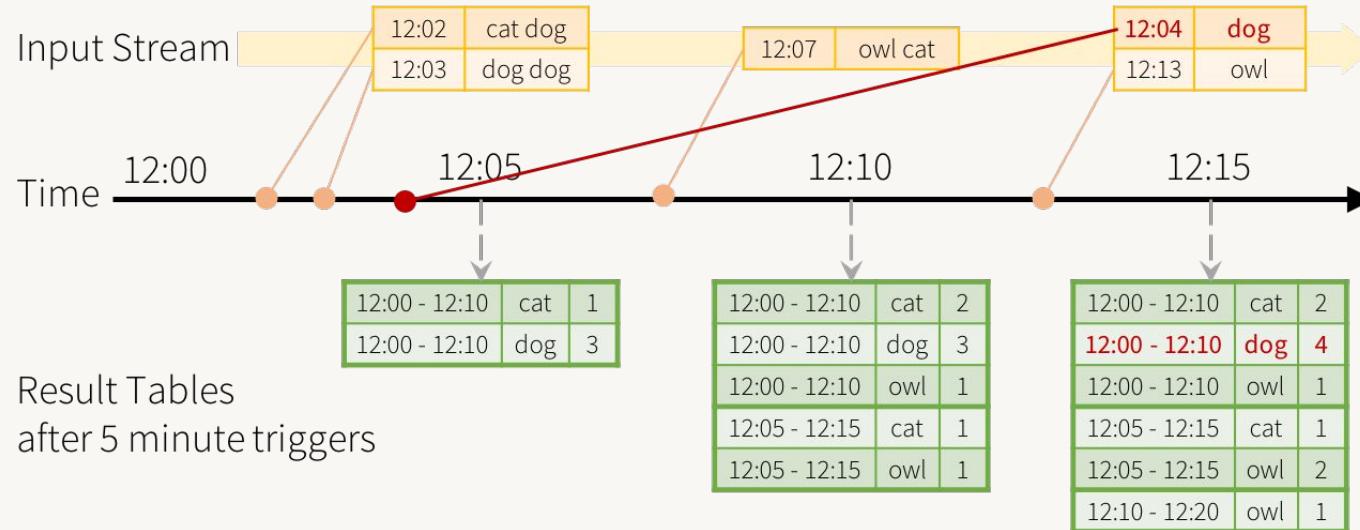
Process based on event-time  
(time fields embedded in data)  
rather than receipt time

## WATERMARKS

Handle late data and limit how  
long to remember old data



# Handling Late Data and Watermarking



# Watermarking

```
(streamingDF
    .withWatermark("time", "2 hours")
    .groupBy(col("device"),
              window(col("time"), "1 hour"))
    .count()
)
```





Module 6

# Delta Lake





# Delta Lake

## Using Spark with Delta Lake

---

# Delta Lake

Delta Lake Concepts

95



# What is Delta Lake?

- Technology designed to be used with Apache Spark to build robust data lakes
- Open source project at [delta.io](https://delta.io)
- Databricks [Delta Lake documentation](#)



**DELTA LAKE**



# Delta Lake features

- ACID transactions on Spark
- Scalable metadata handling
- Streaming and batch unification
- Schema enforcement
- Time travel
- Upserts and deletes
- Fully configurable/optimizable
- Structured streaming support



# Delta Lake components

Delta Lake  
storage  
layer

Delta tables

Delta Engine



# Delta Lake Storage Layer

- Highly performant and persistent
- Low-cost, easily scalable object storage
- Ensures consistency
- Allows for flexibility



# Delta tables

- Contain data in Parquet files that are kept in object storage
- Keep transaction logs in object storage
- Can be registered in a metastore (optional)



# Delta Engine

Databricks edge feature; not available in OS Apache Spark

- File management optimizations
- Auto-optimized writes
- Performance optimization via Delta caching



# What is the Delta transaction log?

- Ordered record of the transactions performed on a Delta table
- Single source of truth for that table
- Mechanism that the Delta Engine uses to guarantee atomicity



# How does the transaction log work?

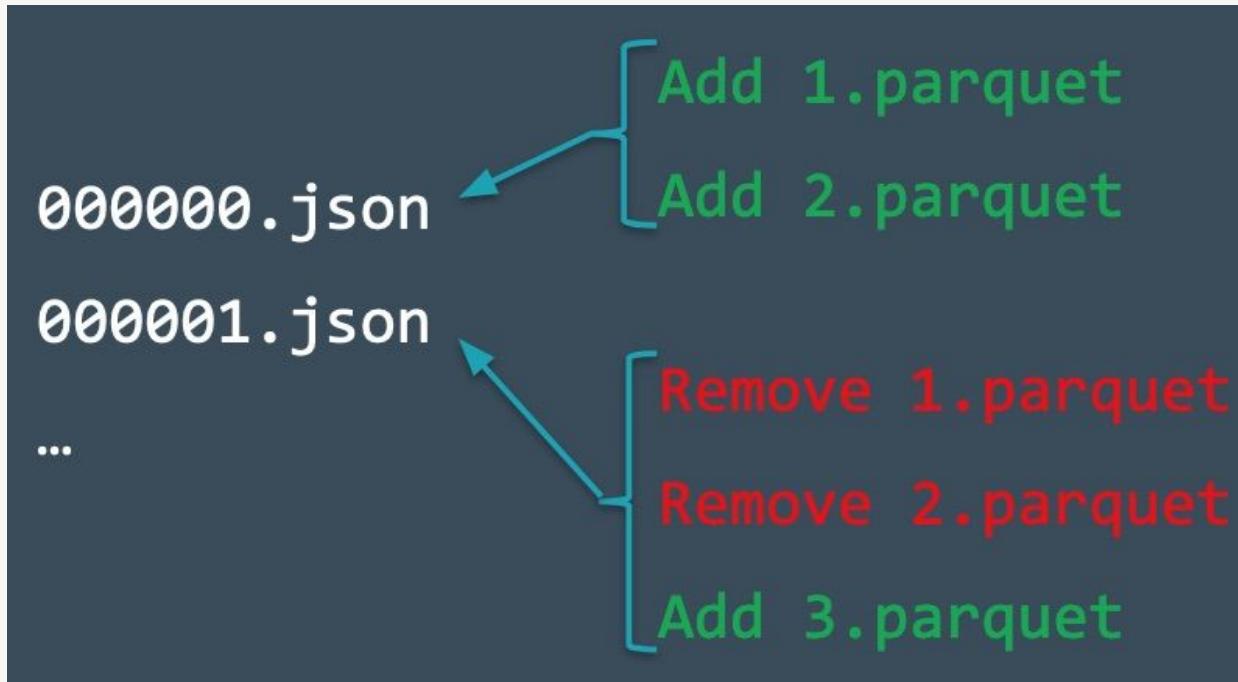
- Delta Lake breaks operations down into one or more of these steps:
  - Add file
  - Remove file
  - Update metadata
  - Set transaction
  - Change protocol
  - Commit info



# Delta transaction log at the file level



# Adding commits to the transaction log



## DESCRIBE Command

- Returns the metadata of an existing table
  - Ex. Column names, data types, comments

## DESCRIBE HISTORY Command

- Returns a more complete set of metadata for a Delta table
  - Operation, user, operation metrics



# Delta Lake Time Travel

- Query an older snapshot of a Delta table
  - Re-creating analysis, reports or outputs
  - Writing complex temporal queries
  - Fixing mistakes in data
  - Providing snapshot isolation



# Time Travel SQL syntax

```
SELECT * FROM events TIMESTAMP AS OF timestamp_expression  
SELECT * FROM events VERSION AS OF version
```



# Time Travel SQL syntax

```
SELECT * FROM events TIMESTAMP AS OF timestamp_expression
```

- `timestamp_expression` can be:
  - String that can be cast to a timestamp
  - String that can be cast to a date
  - Explicit timestamp type (the result of casting)
  - Simple time or date expressions (see the [Databricks documentation](#))



# Time Travel SQL syntax

```
SELECT * FROM events VERSION AS OF version
```

- Version can be obtained from the output of DESCRIBE HISTORY events.



# Thank you! Congrats!

